CS 234 Winter 2019
Assignment 1
Due: January 23 at 11:59 pm

For submission instructions please refer to website

1 Optimal Policy for Simple MDP [20 pts]

Consider the simple n-state MDP shown in Figure 1. Starting from state s_1, the agent can move to the right (a_0) or left (a_1) from any state s_i. Actions are deterministic and always succeed (e.g. going left from state s_2 goes to state s_1, and going left from state s_1 transitions to itself). Rewards are given upon taking an action from the state. Taking any action from the goal state G earns a reward of $r = +1$ and the agent stays in state G. Otherwise, each move has zero reward ($r = 0$). Assume a discount factor $\gamma < 1$.

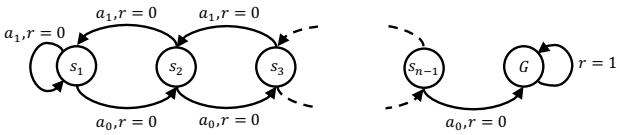

Figure 1: n-state MDP

(a) The optimal action from any state s_i is taking a_0 (right) until the agent reaches the goal state G. Find the optimal value function for all states s_i and the goal state G. [5 pts]

$$
\begin{aligned}
V(s_1) &= R(s_1, a_0) + \gamma V(s_2) \\
V(s_2) &= R(s_2, a_0) + \gamma V(s_3) \\
&\ldots\ldots \\
V(s_{n-1}) &= R(s_{n-1}, a_0) + \gamma V(G) \\
V(G) &= 1 + \gamma V(G)
\end{aligned}
$$

so we can get:

$$V(G) = \frac{1}{1-\gamma}$$
$$V(s_{n-1}) = \frac{\gamma}{1-\gamma}$$
$$V(s_{n-2}) = \frac{\gamma^2}{1-\gamma}$$
$$.....$$
$$V(s_1) = \frac{\gamma^{n-1}}{1-\gamma}$$

(b) Does the optimal policy depend on the value of the discount factor γ? Explain your answer. [5 pts]

Consider the policy in state s_1:

$$Q(s_1, a_0) = \gamma V(s_2) = \frac{\gamma^{n-1}}{1-\gamma}$$
$$Q(s_1, a_1) = \gamma V(s_1) = \frac{\gamma^n}{1-\gamma} < Q(s_1, a_0)$$

if $\gamma > 0$, then $Q(s_1, a_0) > Q(s_1, a_1)$, the optimal action is a_0, as in other states. The more a_1 the policy takes, the more discounted value it get.
if $\gamma < 0$, then the value of each state may be greater or less than zero. So the optimal policy depends.
if $\gamma = 0$, every action value in each state presents the same, indicating that the policy is unrelated with the state transition. There may be several optimal policy.

(c) Consider adding a constant c to all rewards (i.e. taking any action from states s_i has reward c and any action from the goal state G has reward $1 + c$). Find the new optimal value function for all states s_i and the goal state G. Does adding a constant reward c change the optimal policy? Explain your answer. [5 pts]

$$V(G) = \frac{1+c}{1-\gamma}$$
$$V(s_{n-1}) = \frac{c+\gamma}{1-\gamma}$$
$$V(s_{n-2}) = \frac{c+\gamma^2}{1-\gamma}$$
$$.....$$
$$V(s_1) = \frac{c+\gamma^{n-1}}{1-\gamma} = \frac{\gamma^{n-1}}{1-\gamma} + \frac{c}{1-\gamma}$$

All the values is shifted by $\frac{c}{1-\gamma}$, so the optimal policy is unchanged.

(d) After adding a constant c to all rewards now consider scaling all the rewards by a constant a (i.e. $r_{new} = a(c + r_{old})$). Find the new optimal value function for all states s_i and the goal state G. Does that change the optimal policy? Explain your answer, If yes, give an example of a and c that changes the optimal policy. [5 pts]

$$V(G) = a\frac{1+c}{1-\gamma}$$
$$V(s_{n-1}) = a\frac{c+\gamma}{1-\gamma}$$
$$V(s_{n-2}) = a\frac{c+\gamma^2}{1-\gamma}$$
$$\ldots\ldots$$
$$V(s_1) = a\frac{c+\gamma^{n-1}}{1-\gamma} = a\frac{\gamma^{n-1}}{1-\gamma} + a\frac{c}{1-\gamma}$$

if $a > 0$, then the optimal policy is unchanges.
if $a = 0$, then any policy is optimal.
if $a < 0$, then the value increases as the action a_1 is taken, converging to $\frac{ac}{1-\gamma}$, which is greater than $V(G)$. So any policy which never reaches terminal state is the optimal policy.

2 Running Time of Value Iteration [20 pts]

In this problem we construct an example to bound the number of steps it will take to find the optimal policy using value iteration. Consider the infinite MDP with discount factor $\gamma < 1$ illustrated in Figure 2. It consists of 3 states, and rewards are given upon taking an action from the state. From state s_0, action a_1 has zero immediate reward and causes a deterministic transition to state s_1 where there is reward $+1$ for every time step afterwards (regardless of action). From state s_0, action a_2 causes a deterministic transition to state s_2 with immediate reward of $\gamma^2/(1-\gamma)$ but state s_2 has zero reward for every time step afterwards (regardless of action).

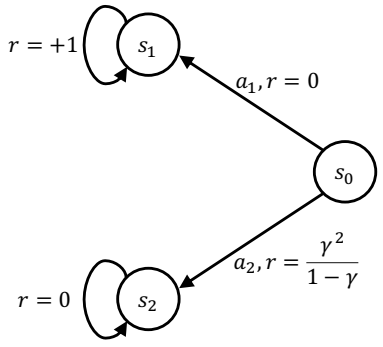

Figure 2: infinite 3-state MDP

(a) What is the total discounted return ($\sum_{t=0}^{\infty} \gamma^t r_t$) of taking action a_1 from state s_0 at time step

$t = 0$? [5 pts]

$$G(s_0) = \sum_{t=1}^{\infty} \gamma^t r_t = \lim_{t \to \infty} \frac{\gamma - \gamma^t}{1 - \gamma} = \frac{\gamma}{1 - \gamma}$$

(b) What is the total discounted return ($\sum_{t=0}^{\infty} \gamma^t r_t$) of taking action a_2 from state s_0 at time step $t = 0$? What is the optimal action? [5 pts]

$$G(s_0) = \sum_{t=1}^{\infty} \gamma^t r_t = \frac{\gamma^2}{1 - \gamma} + 0 + ... = \frac{\gamma^2}{1 - \gamma}$$

so the optimal action is a_1.

(c) Assume we initialize value of each state to zero, (i.e. at iteration $n = 0$, $\forall s : V_{n=0}(s) = 0$). Show that value iteration continues to choose the sub-optimal action until iteration n^* where,

$$n^* \geq \frac{\log(1 - \gamma)}{\log \gamma} \geq \frac{1}{2} \log(\frac{1}{1 - \gamma}) \frac{1}{1 - \gamma}$$

Thus, value iteration has a running time that grows faster than $1/(1 - \gamma)$. (You just need to show the first inequality) [10 pts]

On each iteration, $V(s_0) = \frac{\gamma^2}{1-\gamma}$, but $V(s_1)$ increases by γ^t.After n steps, $V(s_1) = \frac{\gamma - \gamma^{n+1}}{1-\gamma}$. To let the policy choose a_1, what we only need is

$$\begin{aligned} \frac{\gamma - \gamma^{n+1}}{1 - \gamma} &\geq \frac{\gamma^2}{1 - \gamma} \\ 1 - \gamma^n &\geq \gamma \\ n &\leq \frac{log(1 - \gamma)}{log(\gamma)} \\ ??? \end{aligned}$$

3 Approximating the Optimal Value Function [35 pts]

Consider a finite MDP $M = \langle S, A, T, R, \gamma \rangle$, where S is the state space, A action space, T transition probabilities, R reward function and γ the discount factor. Define Q^* to be the optimal state-action value $Q^*(s, a) = Q_{\pi^*}(s, a)$ where π^* is the optimal policy. Assume we have an estimate $\tilde{Q}$ of Q^*, and $\tilde{Q}$ is bounded by l_∞ norm as follows:

$$||\tilde{Q} - Q^*||_\infty \leq \varepsilon$$

Where $||x||_\infty = max_{s,a}|x(s, a)|$.

Assume that we are following the greedy policy with respect to $\tilde{Q}$, $\pi(s) = argmax_{a \in \mathcal{A}} \tilde{Q}(s, a)$. We want to show that the following holds:

$$V_\pi(s) \geq V^*(s) - \frac{2\varepsilon}{1-\gamma}$$

Where $V_\pi(s)$ is the value function of the greedy policy π and $V^*(s) = max_{a \in A} Q^*(s, a)$ is the optimal value function. This shows that if we compute an approximately optimal state-action value function and then extract the greedy policy for that approximate state-action value function, the resulting policy still does well in the real MDP.

(a) Let π^* be the optimal policy, V^* the optimal value function and as defined above $\pi(s) = argmax_{a \in A} \tilde{Q}(s, a)$. Show the following bound holds for all states $s \in S$. [10 pts]

$$V^*(s) - Q^*(s, \pi(s)) \leq 2\varepsilon$$

$$\begin{aligned}
V^*(s) - Q^*(s, \pi(s)) &= V^*(s) - \tilde{Q}(s, \pi(s)) + \tilde{Q}(s, \pi(s)) - Q^*(s, \pi(s)) \\
&= max_{a \in A} Q^*(s, a) - \tilde{Q}(s, \pi(s)) + \tilde{Q}(s, \pi(s)) - Q^*(s, \pi(s)) \\
&= Q^*(s, \pi^*(s)) - \tilde{Q}(s, \pi(s)) + \tilde{Q}(s, \pi(s)) - Q^*(s, \pi(s)) \\
&\leq Q^*(s, \pi^*(s)) - \tilde{Q}(s, \pi(s)) + \varepsilon \\
&\leq Q^*(s, \pi^*(s)) - \tilde{Q}(s, \pi^*(s)) + \varepsilon \\
&\leq 2\varepsilon
\end{aligned}$$

(b) Using the results of part 1, prove that $V_\pi(s) \geq V^*(s) - \frac{2\varepsilon}{1-\gamma}$. [10 pts]

$$\begin{aligned}
V_\pi(s) &= \tilde{Q}(s, \pi(s)) \\
&\geq Q^*(s, \pi(s)) \\
&\geq V^*(s) - 2\varepsilon \\
&\geq V^*(s) - \frac{2\varepsilon}{1-\gamma}
\end{aligned}$$

Now we show that this bound is tight. Consider the 2-state MDP illustrated in figure 3. State s_1 has two actions, "*stay*" self transition with reward 0 and "*go*" that goes to state s_2 with reward 2ε. State s_2 transitions to itself with reward 2ε for every time step afterwards.

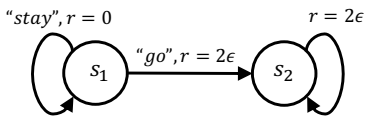

Figure 3: 2-state MDP

(c) Compute the optimal value fucntion $V^*(s)$ for each state and the optimal state-action value function $Q^*(s,a)$ for state s_1 and each action. [5 pts]

$$V^*(s_1) = 2\varepsilon + \sum_{i=1}^{\infty} 2\varepsilon\gamma^i = \frac{2\varepsilon}{1-\gamma}$$
$$V^*(s_2) = V^*(s_1) = \frac{2\varepsilon}{1-\gamma}$$
$$Q^*(s_1, stay) = \tilde{Q}(s_1, stay) = 0 + \gamma V_\pi(s_1) = \frac{2\varepsilon\gamma}{1-\gamma}$$
$$Q^*(s_1, go) = \tilde{Q}(s_1, go) = 2\varepsilon + \gamma V_\pi(s_2) = \frac{2\varepsilon}{1-\gamma}$$

(d) Show that there exists an approximate state-action value function $\tilde{Q}$ with ε error (measured with l_∞ norm), such that $V_\pi(s_1) - V^*(s_1) = -\frac{2\varepsilon}{1-\gamma}$, where $\pi(s) = argmax_{a\in A}\tilde{Q}(s,a)$. (You may need to define a consistent tie break rule) [10 pts]

(c) shows that $V^*(s_1) = \frac{2\varepsilon}{1-\gamma}$. Let the $\pi(s_1) = stay$, then $V_\pi(s_1) = 0$, such that $V_\pi(s_1) - V^*(s_1) = -\frac{2\varepsilon}{1-\gamma}$.

$$\tilde{Q}(s_1, stay) = Q^*(s_1, stay) + \varepsilon$$
$$\tilde{Q}(s_1, go) = Q^*(s_1, go) - \varepsilon$$

4 Frozen Lake MDP [25 pts]

Now you will implement value iteration and policy iteration for the Frozen Lake environment from OpenAI Gym. We have provided custom versions of this environment in the starter code.

(a) **(coding)** Read through `vi_and_pi.py` and implement `policy_evaluation`, `policy_improvement` and `policy_iteration`. The stopping tolerance (defined as $\max_s |V_{old}(s) - V_{new}(s)|$) is tol $= 10^{-3}$. Use $\gamma = 0.9$. Return the optimal value function and the optimal policy. [10pts]

(b) **(coding)** Implement `value_iteration` in `vi_and_pi.py`. The stopping tolerance is tol $= 10^{-3}$. Use $\gamma = 0.9$. Return the optimal value function and the optimal policy. [10 pts]

(c) **(written)** Run both methods on the Deterministic-4x4-FrozenLake-v0 and Stochastic-4x4-FrozenLake-v0 environments. In the second environment, the dynamics of the world are stochastic. How does stochasticity affect the number of iterations required, and the resulting policy? [5 pts]

The number of iterations increase for both the deterministic and the stochastic environment.